

# ~~Inequality Joins in Pandas~~

# ~~with Pyjanitor~~

*Presented by Samuel Oranyeli*

# *Efficient Range Joins in Pandas*

*Presented by Samuel Oranyeli*

# *About Me*

- *Data Engineer*
- *Open source contributor*



# Contents

- *Equi joins vs non-equi joins*
- *What are range joins*
- *Why are range joins slow?*
- A faster option - *conditional\_join*
- *Performance*

# References

- <https://select.dev/posts/snowflake-range-join-optimization>
- <https://www.vertica.com/blog/what-is-a-range-join-and-why-is-it-so-fastba-p223413/>
- <https://duckdb.org/2022/05/27/iejoin.html>
- <https://www.vertica.com/docs/9.2.x/HTML/Content/Authoring/AnalyzingData/Queries/Joins/RangeJoins.htm>
- [Dathan, B et al.](#)
- <https://db.in.tum.de/~reif/papers/p2535-reif.pdf>
- <https://stackoverflow.com/q/24480031/7175713>
- [Zuhair Khayyat et al.](#)

*Equí joins vs non-equí joins*

- □ ×

```
select
...
from orders
join customers
  on orders.customer_id=customers.id -- example equi-join condition
```

- □ ×

```
select
...
from orders
inner join customers
  on orders.customer_id=customers.id -- equi-join
  and orders.created_at > customers.one_year_anniversary_date -- non-equi join
```

— □ ×

```
orders.merge(customers, on = 'customer_id', how = 'inner')
```

— □ ×

```
(orders
    .merge(
        customers,
        on = 'customer_id',
        how = 'inner')
    .loc[lambda df: df.created_at > df.one_year_anniversary_date]
)
```

*What are range joins?*

# Range Joins

*Point in interval join*

— □ ×

```
SELECT owner_id, COUNT(*)  
FROM clicks JOIN ip_owners  
ON clicks.dest_ip  
BETWEEN ip_start AND ip_end  
GROUP BY owner_id;
```

*Overlap join*

— □ ×

```
SELECT  
    r.id,  
    s.id  
FROM events r  
JOIN events s  
ON r.start <= s.end  
AND r.end >= s.start  
AND r.id <> s.id
```

# *Why are range joins slow?*

- Range joins in pandas are usually computed via a cartesian join.
- For relatively large datasets, this becomes expensive memory wise, and is generally slow.
- `IntervalIndex` is an excellent, performant option, but is limited to scenarios where the intervals do not overlap.

- □ ×

pos

0	2
1	3
2	3
3	12
4	20
5	52
6	10

- □ ×

start end

0	1	3
1	5	11
2	19	22
3	30	39
4	7	25

```
(df1
    .merge(df2, how='cross')
    .loc[lambda df: df.pos.between(df.start, df.end)]
)
```

	pos	start	end
0	2	1	3
5	3	1	3
10	3	1	3
19	12	7	25
22	20	19	22
24	20	7	25
31	10	5	11
34	10	7	25

Return rows where df1.pos  
is between df2.start and df2.end

	df1			df2		
	pos			start	end	
0	2			0	1	3
1	3			1	5	11
2	3			2	19	22
3	12			3	30	39
4	20			4	7	25
5	52					
6	10					

IntervalIndex wont be efficient here,  
since the intervals overlap.

cartesian product  
of df1 and df2

pos	start	end
0	2	1
1	2	11
2	2	22
3	2	39
4	2	25
5	3	3
6	3	11
7	3	22
8	3	39
9	3	25
10	3	3
11	3	11
12	3	22
13	3	39
14	3	25
15	12	3
16	12	11
17	12	22
18	12	39
19	12	25
20	20	3
21	20	11
22	20	22
23	20	39
24	20	25
25	52	3
26	52	11
27	52	22
28	52	39
29	52	25
30	10	3
31	10	11
32	10	22
33	10	39
34	10	25

range condition applied  
after cartesian join

pos	start	end
0	2	1
5	3	1
10	3	1
19	12	25
22	20	22
24	20	25
31	10	11
34	10	25

**THERE MUST BE**



**A BETTER WAY**

[memegenerator.net](http://memegenerator.net)

Source: [there must be a better way gif](#)

*conditional\_join*



```
conditional_join(  
    df,  
    right,  
    *args,  
    use_numba,  
    df_columns,  
    right_columns)
```

*\*args*

*(left\_column, right\_column , comparator)*

*pip install pyjanitor*

*import janitor*

```
(df1  
    .conditional_join(  
        df2,  
        ('pos', 'start', '>='),  
        ('pos', 'end', '<='))  
    )
```

	pos	start	end
0	2	1	3
1	3	1	3
2	3	1	3
3	12	7	25
4	20	7	25
5	20	19	22
6	10	5	11
7	10	7	25

Return rows where df1.pos  
is between df2.start and df2.end

	df1			df2		
	pos	start	end	start	end	
0	2			0	1	3
1	3					
2	3	1	5	5	11	
3	12	2	19	19	22	
4	20	3	30	30	39	
5	52	4	7	7	25	
6	10					

' Sort start in ascending order  
' get furthest positions where  
pos is  $\geq$  start  
' get earliest positions where pos is  $\leq$  end

	start	end	cummax
0	1	3	3
1	5	11	11
4	7	25	25
2	19	22	25
3	30	39	39

reduced search space

	pos	bottom	top
0	2	1	0
1	3	1	0
2	3	1	0
3	12	3	2
4	20	4	2
5	52	4	4
6	10	3	1

top should be less than bottom

	pos	start	end
2	2	1	3
3	3	1	3
12	7	25	
20	19	22	
20	7	25	
10	5	11	
10	7	25	

Do actual comparison

	pos	start	end
0	2	1	3
1	3	1	3
2	3	1	3
3	20	19	22
4	12	7	25
5	20	7	25
6	10	5	11
7	10	7	25

final output

# *Sorting + Binary Search*



# *Performance*



```
In [6]: events.shape  
Out[6]: (29999, 6)
```

```
In [7]: events.head()  
Out[7]:
```

	id	name	audience		start		sponsor		end
0	1	Event 1	1178	2022-11-19	10:00:00		Sponsor 2	2022-11-19	10:15:00
1	2	Event 2	1446	2015-09-27	15:00:00		Sponsor 11	2015-09-27	15:11:00
2	3	Event 3	2261	2019-11-12	18:00:00		Sponsor 10	2019-11-12	18:53:00
3	4	Event 4	1471	2019-12-24	22:00:00		Sponsor 6	2019-12-24	22:11:00
4	5	Event 5	2605	2028-06-20	12:00:00		Sponsor 8	2028-06-20	12:31:00

```
In [22]: (events
...: .conditional_join(
...:     events,
...:     ('start', 'end', '<='),
...:     ('end', 'start', '>='),
...:     ('id', 'id', '!='),
...:     use_numba = False,
...:     df_columns = ['id', 'start', 'end'],
...:     right_columns = ['id', 'start', 'end'])
...: )
```

Out[22]:

		left			right					
		id	start	end	id	start	end			
0	10	1993-11-27	12:00:00	1993-11-27	12:37:00	2345	1993-11-27	10:00:00	1993-11-27	12:00:00
1	15	1993-04-04	16:00:00	1993-04-04	18:00:00	11178	1993-04-04	17:00:00	1993-04-04	17:22:00
2	17	2030-10-25	07:00:00	2030-10-25	07:27:00	19605	2030-10-25	06:00:00	2030-10-25	08:00:00
3	26	2005-10-04	17:00:00	2005-10-04	17:18:00	8218	2005-10-04	17:00:00	2005-10-04	17:27:00
4	35	2024-05-02	15:00:00	2024-05-02	15:35:00	6916	2024-05-02	15:00:00	2024-05-02	15:36:00
		...	...	...	...	...	...			
3697	29966	2000-08-26	11:00:00	2000-08-26	13:00:00	29375	2000-08-26	13:00:00	2000-08-26	13:53:00
3698	29971	2018-05-18	04:00:00	2018-05-18	04:18:00	24173	2018-05-18	04:00:00	2018-05-18	04:36:00
3699	29978	1992-06-07	22:00:00	1992-06-07	22:23:00	981	1992-06-07	22:00:00	1992-06-07	22:30:00
3700	29984	2025-06-05	03:00:00	2025-06-05	03:17:00	19051	2025-06-05	01:00:00	2025-06-05	03:00:00
3701	29995	2016-09-04	14:00:00	2016-09-04	14:32:00	12296	2016-09-04	14:00:00	2016-09-04	14:50:00

[3702 rows x 6 columns]

In [23]: %timeit

```
...: (events
...: .conditional_join(
...:     events,
...:     ('start', 'end', '<='),
...:     ('end', 'start', '>='),
...:     ('id', 'id', '!='),
...:     use_numba = False,
...:     df_columns = ['id', 'start', 'end'],
...:     right_columns = ['id', 'start', 'end'])
...: )
```

18.7 ms ± 719 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)

```
In [15]: eventss = events.sample(5000)
```

```
In [16]: (eventss
...: .merge(eventss, how='cross')
...: .loc[lambda df: df.start_x.le(df.end_y) & df.end_x.ge(df.start_y),
...:       ['id_x','start_x', 'end_x', 'id_y','start_y','end_y']]
...: )
```

```
Out[16]:
```

	id_x	start_x	end_x	id_y	start_y	end_y
0	4479	2005-04-26 03:00:00	2005-04-26 03:42:00	4479	2005-04-26 03:00:00	2005-04-26 03:42:00
5001	22513	2012-12-17 20:00:00	2012-12-17 20:44:00	22513	2012-12-17 20:00:00	2012-12-17 20:44:00
10002	27926	2003-12-25 04:00:00	2003-12-25 04:41:00	27926	2003-12-25 04:00:00	2003-12-25 04:41:00
15003	12258	2010-07-15 18:00:00	2010-07-15 18:26:00	12258	2010-07-15 18:00:00	2010-07-15 18:26:00
20004	16255	2013-09-26 07:00:00	2013-09-26 09:00:00	16255	2013-09-26 07:00:00	2013-09-26 09:00:00
...	...	...	...	...	...	...
24979995	29724	1995-11-06 19:00:00	1995-11-06 19:07:00	29724	1995-11-06 19:00:00	1995-11-06 19:07:00
24984996	7677	2028-02-20 01:00:00	2028-02-20 01:47:00	7677	2028-02-20 01:00:00	2028-02-20 01:47:00
24989997	439	2001-04-14 13:00:00	2001-04-14 13:10:00	439	2001-04-14 13:00:00	2001-04-14 13:10:00
24994998	22785	2011-06-23 16:00:00	2011-06-23 18:00:00	22785	2011-06-23 16:00:00	2011-06-23 18:00:00
24999999	28536	2031-10-13 02:00:00	2031-10-13 02:50:00	28536	2031-10-13 02:00:00	2031-10-13 02:50:00

```
[5082 rows x 6 columns]
```

```
In [17]: %%timeit
....: (eventss
....: .merge(eventss, how='cross')
....: .loc[lambda df: df.start_x.le(df.end_y) & df.end_x.ge(df.start_y),
....:       ['id_x','start_x', 'end_x', 'id_y','start_y','end_y']]
....: )
....:
....:
2.07 s ± 101 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)
```

*Performance*



```
In [3]: flights.shape  
Out[3]: (127426, 6)
```

```
In [4]: flights.head()  
Out[4]:
```

	takeoff	landing	orig	dest	start	end
0	2021-11-27 07:15:00	2021-11-27 08:55:00	0	0	2021-11-27 09:40:00	2021-11-27 11:55:00
1	2021-11-27 20:05:00	2021-11-28 00:50:00	1	1	2021-11-28 01:35:00	2021-11-28 03:50:00
2	2021-11-27 21:00:00	2021-11-27 21:35:00	2	2	2021-11-27 22:20:00	2021-11-28 00:35:00
3	2021-11-27 21:15:00	2021-11-27 22:25:00	0	2	2021-11-27 23:10:00	2021-11-28 01:25:00
4	2021-11-26 11:40:00	2021-11-26 14:45:00	3	3	2021-11-26 15:30:00	2021-11-26 17:45:00

- □ ×

```
%%timeit
# classic pandas merge and filter
outt = (flights
        .merge(flights, left_on='dest', right_on='orig')
        .loc[lambda f: f.takeoff_y.between(f.start_x, f.end_x) & (f.orig_x != f.orig_y),
             ['start_x', 'takeoff_y', 'start_y', 'dest_x', 'orig_y']]
        )
8.24 s ± 304 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)
```

```
%%timeit
# use indices from Pandas internal merge function
# compute the non-equi joins
# then create the dataframe
outer = (flights
    .conditional_join(
        flights,
        ('end', 'takeoff', '>='),
        ('start', 'takeoff', '<='),
        ('orig','orig','!='),
        ('dest', 'orig', '=='),
        df_columns = ['start', 'end', 'dest'],
        right_columns = ['takeoff', 'orig'],
        use_numba=False)
)
1.24 s ± 6.72 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)
```

```
%%timeit
# there can be scenarios
# where executing the non-equi joins
# before the equi-join is more performant
out = (flights
    .conditional_join(
        flights,
        ('end', 'takeoff', '>='),
        ('start', 'takeoff', '<='),
        ('orig', 'orig', '!='),
        ('dest', 'orig', '=='),
        df_columns = ['start', 'end', 'dest'],
        right_columns = ['takeoff', 'orig'],
        force=True,
        use_numba=False)
)
```

329 ms ± 1.2 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)

```
%%timeit
# works great if the keys in the equi-join
# are heavily duplicated
# this numba version uses sorting+binary search
# instead of a hash join
out = (flights
       .conditional_join(
           flights,
           ('end', 'takeoff', '>='),
           ('start', 'takeoff', '<='),
           ('orig', 'orig', '!='),
           ('dest', 'orig', '=='),
           df_columns = ['start', 'end', 'dest'],
           right_columns = ['takeoff', 'orig'],
           use_numba=True)
      )
65.9 ms ± 2.65 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)
```

- □ ×

```
# duplicated keys
```

```
In [20]: flights.value_counts(subset='orig')
```

```
Out[20]:
```

```
orig
```

```
3      6257  
0      5452  
25     3535  
44     3332  
54     2847
```

```
...
```

```
178     15  
148     15  
195     15  
153     13  
51      11
```

```
Name: count, Length: 201, dtype: int64
```

```
In [21]: flights.value_counts(subset='dest')
```

```
Out[21]:
```

```
dest
```

```
2      6778  
27     5426  
31     3286  
16     3189  
37     2882
```

```
...
```

```
123     14  
107     12  
115     11  
199     11  
196      7
```

```
Name: count, Length: 201, dtype: int64
```

